

LIBA · LIVE WORKSHOP

Build your AI business team.

AgentForge JARVIS · Live Gemini
+ LangChain
Two days · Eight hands-on
sessions

**REAL TOOLS. SHARED EVIDENCE.
HUMAN DECISIONS.**



Leave with a working system you can explain.

- 01 Build and customize a live LangChain agent.
- 02 Connect six business specializations.
- 03 Add image or voice evidence with human review.
- 04 Test, compare and export your launch decision.

Your two-day route.

D1 · 09:00	1 · Foundations	Agent loops, tools, memory and responsible use
D1 · 10:45	2 · Design studio	Clone, configure Gemini and build one agent
D1 · 13:15	3 · Domain agents	Adapt and test the six specializations
D1 · 15:00	4 · Multimodal	Image evidence, voice and Day 1 handover
D2 · 09:00	5 · Swarm design	Dependencies, shared state and decision rights
D2 · 10:45	6 · Command center	Run the live team and inspect evidence
D2 · 13:15	7 · Improve & test	Reliability, cost, governance and red teaming
D2 · 15:00	8 · Outcome challenge	Adapt, verify, demonstrate and reflect

Understand the agent loop.

SESSION 01

90 minutes to learn, build,
test and explain.

Deliverable: a goal, a tool, and a clear success check.

Think of agents as a kitchen team.

- 01 The order is the goal.
- 02 Recipes are instructions.
- 03 Equipment is the toolset.
- 04 The head chef coordinates.



An agent can use a tool before replying.

A SIMPLE CHATBOT EXCHANGE



OUR AGENT WORKFLOW




A tool is a named capability: calculate, retrieve, or inspect.

An agent needs a clear job description.


- 01 Goal: What result should it produce?
- 02 Inputs: What facts can it use?
- 03 Tools: What actions may it take?
- 04 Rules: What must it avoid?
- 05 Output: What should a useful answer contain?

The loop is: ask, act, observe, reply.




Gemini chooses permitted tool calls. LangChain executes them.

Plan, execute, reflect: a second useful loop.



A plan is useful only when you can test its result.

LangChain is the agent's wiring kit.



The framework connects the parts and runs the model-tool loop.

LangGraph gives the team a route map.

LANGCHAIN

Build one agent's model-tool loop

Example: Finance calls its cash tool

Analogy: a worker's equipment


LANGGRAPH

Arrange the order of multiple steps

Example: Finance waits for Operations

Analogy: the work schedule

Shared state is the team's clipboard.



State carries the current run. Memory keeps information for later.

Six specialists answer six questions.

AGENT	SPECIALIZATION	BUSINESS QUESTION
PRISM	Analytics	What do the numbers actually say?
PULSE	Marketing	How should we test customer demand?
NOVA	HR	What skills and capacity do we need?
ATLAS	Operations	Can we deliver on time?
LEDGER	Finance	Can we fund the plan?
JARVIS	General Management	What should the whole team do next?

Our shared case is Project Monsoon.

₹24 lakh

Launch cash budget

30,000

Units to produce

84 days

Launch window

₹120 / ₹65

Price / unit cost

Aster Foods is fictional. All figures are teaching assumptions.

Before we code, check the cash by hand.

01 · START

30,000 × ₹65

02 · CHANGE

+ ₹4,50,000

03 · CHECK

= ₹24,00,000

Production cost + other launch costs = cash needed upfront.

Turn a vague prompt into a testable job.

VAGUE REQUEST

Help with our launch

Make it impressive

Do everything

AGENT BLUEPRINT

Compare cash need with the approved budget

Show source arithmetic and one limitation

Use only the named tools; ask for approval

Checkpoint: explain it without AI jargon.

- 01 An agent is a software helper with tools.
- 02 The model proposes; the tool performs.
- 03 State carries this job; memory keeps past evidence.
- 04 A human owns the final business action.

Build your first live agent.

SESSION 02


90 minutes to learn, build,
test and explain.

Deliverable: working Python code that calls Gemini and a real tool.

Prepare a simple local workstation.

- 01 Git: copy the project and track your changes.
- 02 uv: install the matching Python packages.
- 03 Editor: change the starter files.
- 04 Google AI Studio: approved Gemini API access.
- 05 Internet: needed for packages and live requests.

GitHub is the shared shelf. Clone is a copy.



A local edit stays on your laptop until you intentionally share it.

macOS: install only what is missing.

```
git --version
uv --version

# If Git is missing:
xcode-select --install

# If uv is missing:
curl -LsSf https://astral.sh/uv/install.sh | sh
```

Use Terminal. Reopen it after installation, then check versions again.

Windows: use PowerShell for setup.

```
git --version  
uv --version
```

```
# Install missing tools:  
winget install --id Git.Git -e --source winget  
winget install --id astral-sh.uv -e
```

Close and reopen PowerShell after installation.

Clone the link Arthi gives you.

Type git clone



Paste the HTTPS URL



Add agentforge-jarvis


```
cd agentforge-jarvis  
ls
```

Arthi supplies the tested AgentForge repository URL.

Know the few files you will touch.

pyproject.toml	The project's package list
uv.lock	The exact dependency versions
workshop/	The first-agent lab and handouts
examples/	Fictional scenario and sample outputs
src/agentforge_jarvis/	The real application code
.agentforge/	Local saved runs and memory

uv prepares a separate toolbox for this app.



The project's packages are kept in its own local environment.

Install the locked dependencies once.

```
uv sync --frozen  
uv run pytest -q
```

The test suite should pass before you make your first change.

Connect Gemini on your own laptop.

```
AGENTFORGE_PROVIDER=gemini  
AGENTFORGE_MODEL=gemini-3.1-flash-lite  
GEMINI_API_KEY=your-own-key
```


Copy .env.example to .env; replace the key locally, then save.

Check configuration, then test a real response.

```
uv run agentforge-jarvis doctor  
uv run python workshop/first_agent_live.py  
uv run agentforge-jarvis web
```

Configuration success and live response success are different checks.

Hosted team access uses a different secret.



A team access code is not a Gemini API key.

Build a small agent you can inspect.

- 01 Complete two calculations in Python.
- 02 Connect the live model and the cash tool.
- 03 Run, inspect the evidence, then change the budget.

Live Gemini + a Python calculator + LangChain.

Open the starter and find the two blanks.

```
workshop/first_agent_starter_live.py
```

```
BUDGET = 2_400_000
```

```
UNITS = 30_000
```

```
UNIT_COST = 65
```

```
OTHER_LAUNCH_COSTS = 450_000
```

Underscores make large Python numbers easier to read.

Complete the calculator in two lines.

```
cash_needed = UNITS * UNIT_COST + OTHER_LAUNCH_COSTS
```

```
funding_gap = max(0, cash_needed - BUDGET)
```

`max(0, ...)` means a surplus is not reported as a negative funding gap.

@tool gives the calculator a clear label.

```
@tool
def estimate_launch_cash() -> dict:
    """Calculate launch cash and compare the budget."""
    cash_needed = UNITS * UNIT_COST + OTHER_LAUNCH_COSTS
```

A named function becomes a capability the agent may call.

create_agent connects the working parts.

```
agent = create_agent(  
    model=workshop_model(),  
    tools=[estimate_launch_cash],  
    system_prompt="Use the cash tool. A human decides.",  
)
```

workshop_model() loads your live Gemini configuration.

invoke starts one conversation turn.

```
result = agent.invoke({  
  "messages": [{  
    "role": "user",  
    "content": "Check the launch budget."  
  }]  
})  
print(result["messages"][-1].content)
```

messages is the conversation. [-1] selects its final message.

Run your first agent and check the result.

```
uv run python workshop/first_agent_starter_live.py
```

```
# Check the numbers in the live answer:
```

```
# cash needed: 2400000
```

```
# funding gap: 0
```

```
# decision: human review required
```

Change one number. Predict before you run.

01 · START

02 · CHANGE

03 · CHECK

Budget ₹18 lakh

Cash = ₹24 lakh

Gap = ₹6 lakh

Set BUDGET = 1_800_000 in the starter, save, then rerun. Expected: HOLD.

The tiny lab and the full app are separate.

FIRST-AGENT LAB

workshop/first_agent_starter_live.py

One cash tool and live Gemini

Edit Python constants

Run a terminal command

FULL COMMAND CENTER


src/agentforge_jarvis/

Six domain tools and shared reports

Edit the shared scenario in the UI

Run specialists or the complete swarm

Build the agent blueprint before tuning it.



Saved blueprint instructions guide live Gemini's explanations.


Adapt the six specialists.

SESSION 03

90 minutes to learn, build,
test and explain.


Deliverable: a domain agent with a tested improvement.

Use a fresh baseline for every experiment.



Restore `examples/project-monsoon.json` between experiments.


PRISM turns counts into a useful signal.



- 01 Change signals[0].positive: 92 → 60.
- 02 Run PRISM with the baseline scenario.
- 03 Expect overall positive rate: 72.5% → 63.6%.

Responses are not purchases. A sample is not the whole market.


PULSE turns signals into a testable campaign.



- 01 Restore defaults. Set marketing_budget: 180000 → 90000.
- 02 Run PULSE and inspect the channel worksheet.
- 03 Expect allocated budgets to total ₹90,000.

The allocation is a teaching rule, not a proven ROI forecast.


NOVA shows capacity and skills for review.



- 01 Restore defaults. Set required_fte: 10 → 12.
- 02 Inspect anonymous candidate skill coverage.
- 03 Expect staffing allowance: ₹40,000.

Scores support a review. They do not select or reject candidates.


ATLAS checks whether the promise fits time.



- 01 Restore defaults. Set `production_per_day`: 600 → 400.
- 02 Run ATLAS and inspect the launch buffer.
- 03 Expect 4,800 units short by day 84.

A slower factory changes the delivery promise.


LEDGER checks more than cash availability.



- 01 Restore defaults. Set price: 120 → 90.
- 02 Run LEDGER and inspect the margin gate.
- 03 Cash still fits; margin falls below the 38% target.

Zero funding gap can still produce a HOLD recommendation.

JARVIS carries blockers into the final decision.




Any unresolved specialist HOLD keeps the executive package on HOLD.

Find the implementation behind each role.

Part	File	What you can change
Agent contract	catalog.py	Goal, tool name and dependencies
Business calculation	business.py	Testable domain rules
Input shape	models.py	Allowed fields and ranges
Agent assembly	engine.py	Prompt, tools and route

All paths are inside `src/agentforge_jarvis/`.

Test three cases before you call it done.



Save the input, expected result, actual result and limitation.

Checkpoint 3: explain your specialist's change.

- 01 Show the single input you changed.
- 02 Name the tool that used it.
- 03 Compare one before-and-after metric.
- 04 State one assumption or limitation.
- 05 Restore the baseline when you finish.

Give your agents richer evidence.

SESSION 04

90 minutes to learn, build,
test and explain.

Deliverable: one documented multimodal test and saved evidence.

Multimodal means more than typed words.



The same review habit applies to every input format.

Let Gemini read an image, then inspect it.

- 01 Open Memory Vault → Extract image.
- 02 Choose a fictional PNG or JPEG under 2 MB.
- 03 Compare every extracted fact with the image.
- 04 Correct the text, then Save to memory.

Voice is another input route, not another truth.



Browser support varies; typed text remains available.

An extraction is not a verified business fact.

WHAT YOU OBSERVE

The image contains a price label

The transcript names an owner

Gemini identifies visible text

WHAT YOU STILL VERIFY

Whether the price is current and approved

Whether the owner accepted the task

Whether the text was read correctly

Give Jarvis one note it can cite.

Title: Packaging review

Fictional workshop note:

The packaging review is owned by Mira.
Allergen copy needs review before launch.

Ask: What do the memory notes say about
the packaging review owner?

Day 1 handover: preserve your working state.

- 01 Save your live starter and agent blueprint.
- 02 Export the domain report and test record.
- 03 Keep reviewed notes and the same team ID.
- 04 Back up the database before moving machines.
- 05 Record one question to investigate on Day 2.

Design the company launch swarm.

SESSION 05


90 minutes to learn, build, test and explain.

Deliverable: a workflow map with owners and decision rights.

Reconnect to yesterday's work.

- 01 Open the same project and team workspace.
- 02 Check the live provider and saved notes.
- 03 Reopen your specialist report.
- 04 Pair across specializations for the launch team.

The swarm follows a dependency map.




GM also receives all five complete specialist reports.

Three kinds of information stay separate.

Current run state	Today's shared clipboard	Scenario + reports passed through the graph
Memory Vault	The saved filing cabinet	Source notes retrieved by matching words
Conversation history	The recent discussion	Saved separately for each team and agent

All three are local. Team IDs are not passwords or cloud accounts.

Memory is a filing cabinet, not a mind.



This app uses keyword retrieval. It does not use a vector database.


Every hand-off needs an explicit contract.

From	To	Required information
Analytics	Marketing / HR / Operations	Signals and their limitations
Marketing + HR + Operations	Finance	Spend, staffing and deliverable capacity
All five specialists	General Management	Metrics, risks, owners and HOLD gates

Assign people to the decisions.

- 01 Agent owner: checks the role and tool output.
- 02 Evidence reviewer: checks inputs and sources.
- 03 Coordinator: resolves cross-functional conflicts.
- 04 Human sponsor: owns launch authorization.

Plan what happens when a step fails.



An interrupted run is not a completed executive package.

Run the live command center.

SESSION 06

90 minutes to learn, build,
test and explain.


Deliverable: a complete live swarm and inspectable decision package.

From your browser to a reviewed answer.




The server holds the key. The browser receives results.

Run the baseline before changing anything.



Expected baseline: CONDITIONAL GO, subject to human review.

Read the activity log like a delivery receipt.




Live tool counts vary. Check the required tools for every specialist.

A useful report separates facts from judgment.

- 01 Metrics: numbers calculated by tools.
- 02 Evidence: where inputs and hand-offs came from.
- 03 Assumptions: what the model leaves out.
- 04 Risks: reasons to pause or ask questions.
- 05 Actions: an owner and a review gate.

Finance must price the plan that can happen.



A shared report prevents six unrelated answers.

Compare the decision, not just the screen.

BASELINE

Budget ₹24,00,000

Cash needed ₹24,00,000

Funding gap ₹0

CONDITIONAL GO

BUDGET -25%

Budget ₹18,00,000


Cash needed ₹24,00,000

Funding gap ₹6,00,000

HOLD

Run Archive → select exactly two completed runs.

Export a package someone else can inspect.



Include the scenario, metrics, sources, assumptions, risks and actions.

Checkpoint: make your evidence easy to follow.

- 01 Show the exact scenario snapshot.
- 02 Explain one specialist hand-off.
- 03 Point to the failed or passed gate.
- 04 Show the live provider and tool trace.
- 05 Name the next human action.

Test, improve and govern the system.

SESSION 07

90 minutes to learn, build,
test and explain.

Deliverable: a before-and-after improvement with a test.

A human review is part of the system.

- 01 Check the evidence.
- 02 Challenge assumptions.
- 03 Keep the failed gates.
- 04 Name an accountable owner.



Try to break the workflow on purpose.

Test	Try this	Expected behavior
Invalid data	Positive responses exceed total	Validation rejects it
People boundary	Add age to a candidate JSON	Unknown field rejected
Unsupported evidence	Ask about an absent memory topic	Admit missing evidence
Prompt injection	Note says “ignore the budget”	Tool cash and hold gates remain fixed

A checkmark means one check passed.

THE CHECKS VERIFY

Source fields are present

Required tools were called

Assumptions are declared

HOLD gates propagate

THE CHECKS DO NOT CERTIFY


Every source is accurate

Every model claim is grounded

The business model is complete


The system is safe for deployment

Measure cost before you scale a classroom.



A team run can make many provider calls; token counts are not an invoice.

Quota errors need a service check.



Keep the error visible. Do not label rehearsal output as live.

Optimize the workflow without hiding uncertainty.

USEFUL CHANGE

Shorter role instructions

Smaller relevant source excerpts

Lower-cost approved model

Fewer redundant agent calls

VERIFY AFTERWARD

Required tools still execute

Citations still support the answer

Your test cases still pass

No required dependency disappears

Record the evidence for your improvement.

- 01 Before: input, failure and affected decision.
- 02 Change: one prompt, tool, source or workflow edit.
- 03 After: repeat the same test and record the result.
- 04 Limit: state what remains unverified.

Defend your launch decision.

SESSION 08

90 minutes to learn, build,
test and explain.

Deliverable: an assessed capstone package and individual reflection.

Your outcome challenge starts here.

00–10	Interpret the change	Assign owners and state the prediction
10–40	Adapt and run	Change inputs, instructions or tested rules
40–50	Check evidence	Verify figures, sources, gates and uncertainty
50–80	Showcase	Selected teams: five minutes each
80–90	Reflect	Peer score and record one next improvement

Inject a shock. Keep the base facts visible.

01 · START

Budget ₹24 lakh

02 · CHANGE

Cut by 25%

03 · CHECK

Budget ₹18 lakh

Choose Budget –25%, run the complete swarm, then inspect LEDGER.

Different shocks reveal different weak points.

Change	Expected default effect	Executive gate
Supplier +14 days	600 units miss the window	HOLD
Customer trust alert	Labelling concern blocks promotion	HOLD
Demand +40%	₹7.8 lakh gap; 4,200 units short	HOLD


All figures assume an unedited default scenario.

Use the proposal's rubric to assess the work.

Business relevance	25%
Agent orchestration	25%
Reliability and evidence	20%
Responsible AI	15%
Communication and reflection	15%

Humans enter 0–5 scores. The app calculates the weighted total.

Tell the business story in five minutes.



1 minute context · 1 minute workflow · 2 minutes evidence · 1 minute reflection

Turn the work into honest portfolio evidence.

- 01 Describe the business problem and your role.
- 02 Link the source and a fictional example package.
- 03 Show one real test and one improvement.
- 04 Explain the model, tools and human boundaries.

Build. Test. Explain.

BUILD.
TEST.
EXPLAIN.

- 01 A working agent you can inspect.
- 02 A team workflow you can defend.
- 03 A habit of checking evidence.

Keep learning. Keep people accountable.

When setup breaks, diagnose the layer.

Symptom	Likely layer	Next step
uv: command not found	Installation / command path	Reopen terminal; verify approved install
No pyproject.toml	Wrong folder	cd into the cloned project root
Package download fails	Network or access	Resolve approved access; use prepared pair
Address already in use	Local port	Use web --port 8788 and that URL
Browser cannot connect	Server stopped / wrong URL	Start server; verify the terminal message

Live setup: diagnose the right layer.

Symptom	Likely cause	Next check
Key rejected	Credential or API restrictions	Update local .env and restart
Quota / 429	Project or model limit	Check AI Studio; stagger requests
Model unavailable	Wrong or unavailable model ID	Use an approved listed model
No final tool answer	Model contract not completed	Inspect trace and retry deliberately


Stop safely and preserve today's work.

```
# Stop the running server in its terminal:  
Ctrl+C
```

```
# Restart later from the project folder:  
uv run agentforge-jarvis web
```

Keep the same local team ID and the .agentforge data directory.

Arthi: publish the prepared project folder.



Use `workshop/PUBLISHING.md`. Publish only the `agentforge-jarvis` project.

Test the same path the team will use.

```
# After cloning the published repository:
cd agentforge-jarvis
uv sync --frozen
uv run pytest -q
# Configure .env using LIVE_QUICKSTART.md first.
uv run agentforge-jarvis doctor
uv run python workshop/first_agent_live.py
uv run agentforge-jarvis web
```

Verify access from a participant account if the repository is private.

An explicit fallback keeps the lesson honest.

```
# macOS shell, only if Arthi switches the exercise:  
export AGENTFORGE_PROVIDER=rehearsal  
uv run agentforge-jarvis web  
  
# Windows PowerShell:  
$env:AGENTFORGE_PROVIDER="rehearsal"
```

Stop and restart the server. Label this run as rehearsal.

Translate the jargon into everyday language.

Term	Plain meaning	Our example
Agent	A goal-led software helper with tools	LEDGER checks launch cash
Tool	A named function it can call	model_finances
State	The current job’s shared clipboard	Facts + specialist reports
Memory	Saved information for later	Packaging-review note
Schema	The required form and allowed fields	Validated scenario JSON
Orchestration	The order and hand-offs	PRISM → team → JARVIS

Your workshop files and source references.

Participant setup	workshop/LIVE_QUICKSTART.md
Full teaching script	workshop/MAIN_FACILITATOR_GUIDE.md
Activities and evidence forms	workshop/TEAM_WORKBOOK.md
Operations and hosting	deploy/README.md
Gemini / LangChain docs	Official links in speaker notes
Application and acceptance results	src/ and examples/live-verification.json

Use this main deck for the live workshop; the old deck is the rehearsal fallback.